

This is a test document for chunk extraction.
It contains multiple sentences and paragraphs.
The chunk extraction system should be able to split this text
into smaller chunks of approximately 500 words each.
This will help test the semantic chunking functionality.
The system uses word-based splitting with overlap.
Each chunk should contain meaningful text segments.
This is particularly important for RAG applications.
The chunks will be used to provide context to AI models.
Proper chunking improves retrieval accuracy.
This test document demonstrates the basic functionality.
Additional content line 1 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 2 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 3 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 4 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 5 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 6 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 7 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 8 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 9 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 10 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 11 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 12 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 13 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 14 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 15 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 16 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 17 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 18 for testing chunk extraction.
This line contains more text to increase the document length.
Additional content line 19 for testing chunk extraction.
This line contains more text to increase the document length.

This is the second page of the test document.
It demonstrates multi-page document processing.
The chunk extractor should handle multiple pages correctly.
Each page's content should be processed separately.
Then combined into appropriate chunks.
Page boundaries should be respected in chunk metadata.